

```
In [1]: 1 import pandas as pd
        2
        3 df = pd.read_csv('personal_transactions.csv')
        4 df
```

Out[1]:

	Date	Description	Amount	Transaction Type	Category	Account Name
0	01/01/2018	Amazon	11.11	debit	Shopping	Platinum Card
1	01/02/2018	Mortgage Payment	1247.44	debit	Mortgage & Rent	Checking
2	01/02/2018	Thai Restaurant	24.22	debit	Restaurants	Silver Card
3	01/03/2018	Credit Card Payment	2298.09	credit	Credit Card Payment	Platinum Card
4	01/04/2018	Netflix	11.76	debit	Movies & DVDs	Platinum Card
...
801	09/27/2019	Biweekly Paycheck	2250.00	credit	Paycheck	Checking
802	09/28/2019	BP	33.46	debit	Gas & Fuel	Platinum Card
803	09/28/2019	Sheetz	4.27	debit	Gas & Fuel	Platinum Card
804	09/30/2019	Starbucks	1.75	debit	Coffee Shops	Platinum Card
805	09/30/2019	Internet Service Provider	75.00	debit	Internet	Checking

806 rows × 6 columns

```
In [2]: 1 print(df.isnull().sum())
```

```
Date          0
Description    0
Amount         0
Transaction Type 0
Category       0
Account Name   0
dtype: int64
```



```
In [5]: 1 df.describe()
```

Out[5]:

	Amount
count	806.000000
mean	273.391489
std	667.630374
min	1.750000
25%	15.687500
50%	37.480000
75%	117.680000
max	9200.000000

Exploratory Data Analysis

```
In [4]: 1 import seaborn as sns
2 import matplotlib.pyplot as plt
3
4 sns.histplot(df['Amount'], bins=50)
5 plt.title('Distribution of Transaction Amounts')
6 plt.show()
7
8 df.groupby('Category')['Amount'].sum().plot(kind='bar')
9 plt.title('Spending by Category')
10 plt.show()
```



Feature Engineering

```
In [6]: 1 # Convert 'Date' to datetime
2 df['Date'] = pd.to_datetime(df['Date'])
3 df['Year'] = df['Date'].dt.year
4 df['Month'] = df['Date'].dt.month
5
6 # Feature from 'Description'
7 df['Description_length'] = df['Description'].apply(len)
8
9 df = pd.get_dummies(df, columns=['Transaction Type', 'Account Name'], drop_first=True)
```

Building the Machine Learning Model

```
In [8]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.ensemble import RandomForestClassifier
3 from sklearn.metrics import accuracy_score
4
5 X = df[['Amount', 'Description_length', 'Year', 'Month']]
6 y = df['Category']
7
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```


1 Train the model

```
In [9]: 1 model = RandomForestClassifier(n_estimators=100)
2 model.fit(X_train, y_train)
3
4 y_pred = model.predict(X_test)
5 print('Accuracy:', accuracy_score(y_test, y_pred))
```


Accuracy: 0.9008264462809917

Visualizing Spending Patterns (Data Visualization)

```
In [10]: 1 # Visualize spending by month
2 df.groupby('Month')['Amount'].sum().plot(kind='line')
3 plt.title('Spending by Month')
4 plt.show()
```



```
In [31]: 1 #Spending Breakdown by Category
2
3 df.groupby('Category')['Amount'].sum().plot(kind='pie', autopct='%1.1f%%')
4 plt.title('Spending Breakdown by Category')
5 plt.show()
```



Advanced Features (Optional)

Budget Alerts:

Implement a function to check if the total spending exceeds a set budget:

```
In [13]: 1 def check_budget(df, budget):
2     total_spent = df['Amount'].sum()
3     if total_spent > budget:
4         print(f"Alert: You have exceeded your budget by {total_spent - budget}")
5     else:
6         print(f"Budget is on track. Total spent: {total_spent}")
```

```
In [15]: 1 import pandas as pd
2
3 df = pd.read_csv('personal_transactions.csv')
4 df['Date'] = pd.to_datetime(df['Date'])
5 df.set_index('Date', inplace=True)
```

```
In [20]: 1 # Resample data by month and sum amounts
2 monthly_data = df.resample('M').sum(numeric_only=True)
3 print("Columns in monthly_data:", monthly_data.columns)
4
```

Columns in monthly_data: Index(['Amount'], dtype='object')

```
In [21]: 1 print("Missing values in monthly_data:")
2 print(monthly_data.isnull().sum())
```

Missing values in monthly_data:
Amount 0
dtype: int64

```
In [22]: 1 monthly_data = monthly_data.fillna(0)
2
3 monthly_data['Amount'] = pd.to_numeric(monthly_data['Amount'], errors='coerce')
4 data = monthly_data['Amount']
5 print(data.head())
```

Date
2018-01-31 10094.34
2018-02-28 8385.80
2018-03-31 10821.66
2018-04-30 13196.42
2018-05-31 16483.58
Freq: M, Name: Amount, dtype: float64


Time-Series Forecasting with ARIMA

```
In [26]: 1 import pmdarima
2 print(pmdarima.__version__)

2.0.4
```

In [54]:

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 from statsmodels.tsa.arima.model import ARIMA
4 from sklearn.metrics import mean_squared_error
5
6 df['Date'] = pd.to_datetime(df['Date'])
7 df.set_index('Date', inplace=True)
8 df['Amount'] = pd.to_numeric(df['Amount'], errors='coerce')
9 df.dropna(subset=['Amount'], inplace=True)
10 monthly_data = df.resample('M').sum()
11
12 train = monthly_data[:-12]
13 test = monthly_data[-12:]
14
15 model = ARIMA(train['Amount'], order=(5,1,0)) # Adjust (p,d,q) parameters # Build and fit the ARIMA model
16 model_fit = model.fit()
17
18 forecast = model_fit.forecast(steps=12) # Forecast for the test period
19
20 plt.figure(figsize=(10,6)) # Plot actual vs forecasted data
21 plt.plot(train.index, train['Amount'], label='Training Data')
22 plt.plot(test.index, test['Amount'], label='Actual Data')
23 plt.plot(test.index, forecast, label='Forecasted Data', linestyle='--')
24 plt.title('ARIMA Time Series Forecasting')
25 plt.xlabel('Date')
26 plt.ylabel('Amount')
27 plt.legend()
28 plt.show()
29
30 mse = mean_squared_error(test['Amount'], forecast)
31 print(f'Mean Squared Error: {mse}')
```



Mean Squared Error: 70812305.2179404

- 1
- Blue Line (Training Data): This represents the actual historical data used to train the ARIMA model up to a certain point in time (before 2018-09). The training data shows the trend and seasonal patterns up to the cutoff date.
- 2
- 3
- Orange Line (Actual Data): This is the actual data after the training period, used to compare against the forecasted data. It shows the true values beyond the training window.
- 4
- 5
- Green Dashed Line (Forecasted Data): This line represents the forecasted values by the ARIMA model for the period after the training data. It projects the expected trend based on historical data patterns.
- 6
- 7
- The forecasted data deviates from the actual data, indicating that while the model captures general trends, there are differences between what was predicted and what actually occurred during that time frame.
- 8
- 9
- Key Observations:
- 10
- The ARIMA model reasonably forecasts the general trend but with notable inaccuracies at certain points compared to actual data.
- 11
- The forecast begins around late 2018, and the orange line (actual data) shows fluctuations that the green dashed line doesn't fully capture.

Clustering (Unsupervised Learning)


Goal: Group similar transactions together based on features

```
In [49]: 1 from sklearn.preprocessing import StandardScaler
2 from sklearn.cluster import KMeans
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 df['Transaction Type'] = df['Transaction Type'].astype('category').cat.codes      # Convert categorical variable
7 df['Category'] = df['Category'].astype('category').cat.codes
8 df['Account Name'] = df['Account Name'].astype('category').cat.codes
9
10 X = df[['Amount', 'Category', 'Transaction Type', 'Account Name']]                # Select the relevant columns f
11
12 scaler = StandardScaler()
13 X_scaled = scaler.fit_transform(X)
14
15 wcss = []
16 for i in range(1, 11):
17     kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42)
18     kmeans.fit(X_scaled)
19     wcss.append(kmeans.inertia_)
```

	Date	Description	Amount	Transaction Type	\
0	01/01/2018	Amazon	11.11	debit	
1	01/02/2018	Mortgage Payment	1247.44	debit	
2	01/02/2018	Thai Restaurant	24.22	debit	
3	01/03/2018	Credit Card Payment	2298.09	credit	
4	01/04/2018	Netflix	11.76	debit	

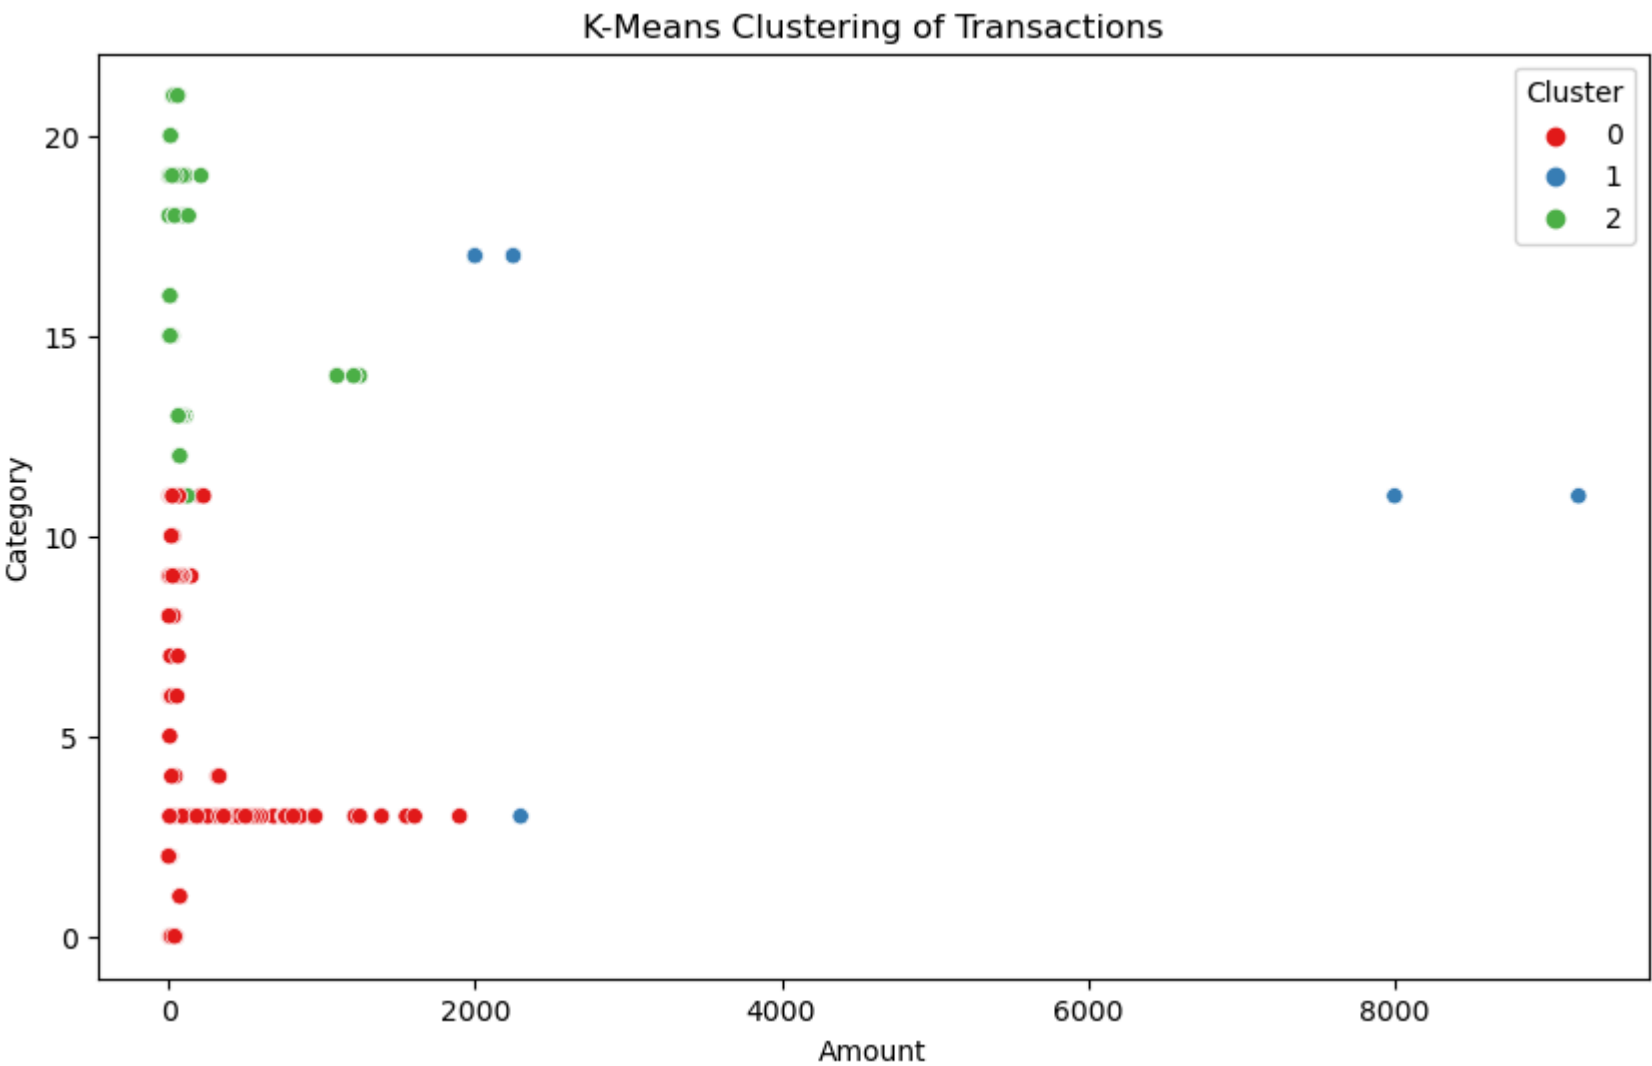
	Category	Account Name
0	Shopping	Platinum Card
1	Mortgage & Rent	Checking
2	Restaurants	Silver Card
3	Credit Card Payment	Platinum Card
4	Movies & DVDs	Platinum Card

```
In [50]: 1 plt.figure(figsize=(8, 5))                                # Plotting the Elbow Curve
2 plt.plot(range(1, 11), wcss, marker='o', linestyle='--')
3 plt.title('Elbow Method to Determine Optimal Number of Clusters')
4 plt.xlabel('Number of Clusters')
5 plt.ylabel('WCSS')      # Within-cluster sum of squares
6 plt.show()
```



```
In [51]: 1 optimal_clusters = 3
2
3 kmeans = KMeans(n_clusters=optimal_clusters, init='k-means++', random_state=42)
4 df['Cluster'] = kmeans.fit_predict(X_scaled)
5 print(df['Cluster'].value_counts())
6
7 # Visualize the clusters (Amount vs Category as an example)
8 plt.figure(figsize=(10, 6))
9 sns.scatterplot(x=df['Amount'], y=df['Category'], hue=df['Cluster'], palette='Set1')
10 plt.title('K-Means Clustering of Transactions')
11 plt.show()
12
13 df.to_csv('clustered_transactions.csv', index=False)
```

```
0    440
2    315
1     51
Name: Cluster, dtype: int64
```



Anomaly Detection

Isolation Forest / LOF: identify unusual or fraudulent transactions. They are effective for detecting outliers in transactional data.

```
In [53]: 1 from sklearn.ensemble import IsolationForest
2
3 features = df[['Amount']] # Add more features if available
4
5 # Initialize and fit the Isolation Forest model
6 iso_forest = IsolationForest(contamination=0.01)
7 df['anomaly'] = iso_forest.fit_predict(features)
8
9 # -1 indicates an anomaly
10 print(df[df['anomaly'] == -1])
```

	Date	Description	Amount	Transaction Type	\
3	01/03/2018	Credit Card Payment	2298.09	credit	
172	05/11/2018	Mike's Construction Co.	8000.00	debit	
676	06/20/2019	Mike's Construction Co.	9200.00	debit	
	Category	Account Name	anomaly		
3	Credit Card Payment	Platinum Card	-1		
172	Home Improvement	Checking	-1		
676	Home Improvement	Checking	-1		

Local Outlier Factor (LOF)

In [54]:

```
1 from sklearn.neighbors import LocalOutlierFactor
2
3 lof = LocalOutlierFactor(n_neighbors=20)
4 df['anomaly'] = lof.fit_predict(features)
5 print(df[df['anomaly'] == -1])
```

	Date	Description	Amount	Transaction Type	\
0	01/01/2018	Amazon	11.11	debit	
3	01/03/2018	Credit Card Payment	2298.09	credit	
11	01/11/2018	Shell	34.87	debit	
28	01/25/2018	Internet Service Provider	69.99	debit	
29	01/29/2018	Shell	30.42	debit	
..	
793	09/18/2019	Credit Card Payment	1606.46	debit	
796	09/20/2019	Credit Card Payment	9.43	credit	
798	09/23/2019	Credit Card Payment	9.43	debit	
801	09/27/2019	Biweekly Paycheck	2250.00	credit	
804	09/30/2019	Starbucks	1.75	debit	

	Category	Account Name	anomaly
0	Shopping	Platinum Card	-1
3	Credit Card Payment	Platinum Card	-1
11	Gas & Fuel	Platinum Card	-1
28	Internet	Checking	-1
29	Gas & Fuel	Silver Card	-1
..
793	Credit Card Payment	Checking	-1
796	Credit Card Payment	Silver Card	-1
798	Credit Card Payment	Checking	-1
801	Paycheck	Checking	-1
804	Coffee Shops	Platinum Card	-1

[134 rows x 7 columns]

Sentiment Analysis on Descriptions

In [58]:

```
1 from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
2
3 analyzer = SentimentIntensityAnalyzer()
4 df['sentiment'] = df['Description'].apply(lambda x: analyzer.polarity_scores(x)['compound'])
5 print(df[['Description', 'sentiment']])
```

	Description	sentiment
0	Amazon	0.1779
1	Mortgage Payment	0.0000
2	Thai Restaurant	0.0000
3	Credit Card Payment	0.3818
4	Netflix	0.0000
..
801	Biweekly Paycheck	0.0000
802	BP	0.0000
803	Sheetz	0.0000
804	Starbucks	0.0000
805	Internet Service Provider	0.0000

[806 rows x 2 columns]